

## Where to Publish SARDI-LoRA — Verified Venue Guide

**Purpose of this document:** Two rounds of AI-assisted research were done to find workshops and student programs for our paper. The first round was honest but vague. The second round confidently invented several workshop names and URLs. **This document keeps only what was independently verified** — every link below was actually checked and works. Anything unverified is clearly marked.

---

### 1. Real, Verified Workshop Matches for Our Project

#### **Best match: DiffuLM — Diffusion Language Models: Foundations, Efficiency, and Reasoning**

- **Venue:** NeurIPS 2026 (Sydney), officially accepted (confirmed on NeurIPS's own accepted-workshops list)
- **Why it fits:** Our entire system *is* a diffusion language model project — this is the most directly matched workshop scope of anything we found.
- **Link:** <https://7amin.github.io/diffulm-neurips2026/>
- **Action:** Open this link and read their actual CFP for exact deadline, page limit, and submission process — the details weren't in the source we found, only the confirmed existence of the workshop.

#### **Strong match: LIGHT — Deployable Small Foundation Models**

- **Venue:** NeurIPS 2026 (Paris), officially accepted
- **Why it fits:** Our project is explicitly about a small, resource-constrained model (172M params) made more capable via LoRA under an 8GB GPU budget — "deployable small foundation models" is almost a direct description of our contribution.
- **Link:** <https://almaai-disi-unibo.github.io/neurips2026-light-smallModels/>
- **Action:** Check their CFP for scope details and deadline.

#### **Possible match: SLM-Agents — 1st Workshop on SLMs for Agentic Systems**

- **Venue:** NeurIPS 2026 (Paris), officially accepted (this confirms the earlier screenshot was right about this one!)
- **Why it fits:** "Small language models" scope matches; less certain fit since "agentic systems" framing may not match our QA/retrieval framing exactly — read their CFP carefully before assuming it fits.
- **Link:** <https://slmw2026.github.io/>

### Also worth checking (less certain fit, but plausible)

- **AXIOM — Foundations of Efficient Deep Learning** (NeurIPS 2026, Paris): <https://axiom-neurips2026.github.io/>
- **On-Device Intelligence: Foundation Models under Real-World Constraints** (NeurIPS 2026, Sydney): <https://odi2026.github.io/>

**None of these deadlines are confirmed yet** — the NeurIPS blog post explicitly said individual workshops are still finalizing their own submission dates (many suggested around Aug 29, 2026, but that's a general target, not a confirmed per-workshop date). **Someone needs to visit each link above and record the actual deadline before we plan our timeline around it.**

---

## 2. Real Student-Specific Programs (Free, Legitimate)

These were separately verified as real, official pages — not invented.

### ACL 2026 / EACL 2026 Student Research Workshop (SRW)

- **What it is:** A dedicated venue for students, with a **free pre-submission mentorship program** — you submit a draft, get anonymously paired with an experienced researcher for feedback *before* the real deadline.
- **Why this matters for us:** Our paper is NLP/QA research — this is a strong scope fit, not a stretch.
- **EACL 2026 SRW:** <https://2026.eacl.org/calls/srw/> (mentions a travel grant application deadline of Feb 10, 2026 — may already be close/passed by the time we're paper-ready; check current status)
- **ACL 2026 SRW:** [https://2026.aclweb.org/calls/student\\_research\\_workshop/](https://2026.aclweb.org/calls/student_research_workshop/)
- **Action:** Read both pages for current deadlines, since ours may fall in a later cycle.

### ML4H (Machine Learning for Health) 2026 — Author Mentorship Program

- **What it is:** Pairs less-experienced authors with senior researchers, free, up to 2 months of mentorship on project direction, experimental design, and writing.
  - **Fit for us:** Weaker — this is a health-ML venue and our project has no health angle. Only worth pursuing if we can't find a better-fit mentorship elsewhere.
  - **Link:** <https://ml4h.ahli.cc/submit/mentorship-programs/>
- 

## 3. Confirmed Policy Facts (from official pages, not guessed)

- **ICLR retired the standalone "Tiny Papers" track after 2024.** Since 2025, individual ICLR workshops are asked to host their own short-paper track (page limits vary by workshop, usually 2–5 pages). Source: <https://iclr.cc/Conferences/2025/CallForTinyPapers>
- **AI-generated papers are explicitly banned from ICLR's tiny/short paper tracks.** The exact wording: *"submissions must be primarily human-authored, reflecting original thought and analysis."* AI **assistance** is allowed, AI **generation** is not. Source: <https://iclr.cc/Conferences/2026/WorkshopGuide>
- **NeurIPS workshop organizers are explicitly told "no hallucinations are allowed"** in their own proposal wording (per NeurIPS's own 2026 workshop guidance) — a good sign the broader NeurIPS community takes this seriously, and a reminder we should hold our own paper to the same standard.

**Practical implication for us:** Since we used AI tools throughout (code debugging, drafting help), we must make sure the *final submitted paper* reflects our own understanding, decisions, and voice — not a lightly-edited AI draft. Disclose AI assistance where required; don't submit AI-generated prose as our own original thought.

---

#### 4. What We Still Don't Know (needs manual verification)

For every workshop link above, someone on the team needs to visit the actual page and record:

- ☐ Exact submission deadline
- ☐ Page limit and required template (NeurIPS/ICLR/ICML style)
- ☐ Whether double-blind/anonymized submission is required
- ☐ Dual submission policy (can we submit to more than one workshop?)
- ☐ Any student travel grant or fee waiver, and its application process/deadline

None of this should be assumed or estimated — copy it directly from the real page.

---

#### 5. Recommended Next Steps, in Order

1. **This week:** Assign one person to open all 5 verified workshop links above and fill in the checklist in Section 4 for each.
  2. **In parallel:** Check EACL/ACL SRW mentorship deadlines — if the mentorship window is still open, submitting our draft there is low-risk, free, and directly improves the paper before any real deadline.
  3. **Pick 1 primary + 1 backup workshop** once deadlines are known — DiffuLM and LIGHT are the strongest scope matches found so far.
  4. **Do NOT** search for or rely on any workshop name that isn't in this document or independently re-verified with a working link.
  5. Once we have a real deadline, work backward to set our own internal draft-complete date, leaving time for the SRW mentorship round if we use it.
- 

*Compiled from AI-assisted research, cross-checked by directly visiting each source URL. Two links from an earlier research pass were found to be fabricated and are intentionally excluded here.*